

MitchAI

Hospitality operations and API-connected n8n workflows

Project overview · Offline portfolio work · reviewed 14 September 2026

MitchAI · recorded workflow checks

Saved synthetic n8n run · 13 September 2026 · historical evidence

Workflow route	HTTP	Expected
daily-operations	200	200
menu-innovation	200	200
social-content	200	200
business-strategy	200	200
customer-feedback	200	200
ai-request	200	200
parallel-ai	200	200
mitch-master · operations	200	200
mitch-master · combo plan	200	200
daily-operations · failure scenario	500	500

10 saved scenarios · all observed HTTP statuses match expected values

Model endpoints and database sink substituted · no live-provider claim

Problem

Hospitality teams repeat daily operations, menu, content and feedback tasks. I used my operating experience to model these as inspectable workflows with explicit inputs and responses.

Engineering decision

A TypeScript/Node.js platform connects eight n8n workflows. Synthetic HTTP scenarios check request handling, provider failure and parallel-branch merging before connecting external systems.

My contribution

Founder and engineer: hospitality requirements, backend implementation, workflow integration and test documentation.

Visual evidence

Visual summary of saved synthetic execution responses from 13 September 2026. Expected HTTP 500 checks a failure path; model endpoints and database sink are substituted.

Engineering evidence

Dated verification, implementation scope and remaining work

Implementation

The backend uses Node.js, TypeScript and Express; the frontend uses Next.js/React with PostgreSQL persistence. Supporting configuration includes Redis and a Qdrant service in Docker Compose. Provider routing includes configurable cloud providers and local Ollama/LM Studio integrations.

The n8n master routes to menu, social, strategy, feedback and operations workflows. The combo route currently returns a plan without executing its steps.

Dated execution evidence

The 13 September verification recorded 193 passing backend tests with none skipped. Eight n8n workflows were exercised through ten synthetic HTTP scenarios, including provider failure and parallel-branch merging.

Page 1 copies route names and observed/expected HTTP statuses from saved responses. All ten match their expected status, including the deliberate 500. Model endpoints and database persistence were substituted. This figure is not a new run or evidence of live model quality or database durability.

Status and evidence access

Offline portfolio prototype. The source repository was made private on 14 September 2026 after historical privacy and credential findings. Anonymous access checks now return 404; credential revocation and history cleanup remain pending. This case study provides the reviewed evidence.

A historical GitHub Actions run on 12 February 2026 built an image and updated AWS ECS. Published Terraform modules are separate implementation evidence, not proof that they provisioned that run. No current service health, SLA, customer usage or independently assessed compliance is claimed.

Next validation

Validate real-provider behaviour and database persistence, resolve the public-history findings and complete security/privacy review before external use. Human approval remains the intended boundary for external posts, messages, filings and purchases.